

1.

s	a	s'	p(s' s,a)	r(s,a,s')	States indexes indicate the following: Coin 1 Value, Coin 2 Value, Where Arm is Over
H,H,1	Call Fairy	H,H,1	1	20	
H,H,2	Call Fairy	H,H,2	1	20	
H,H,1	Move	H,H,2	1	-2	
H,H,2	Move	H,H,1	1	-2	
H,H,1	Flip	H,H,1	0.5	-1	
H,H,1	Flip	T,H,1	0.5	-1	
H,H,2	Flip	H,T,2	0.5	-1	
H,H,2	Flip	H,H,2	0.5	-1	
T,T,1	Call Fairy	T,T,1	1	10	
T,T,2	Call Fairy	T,T,2	1	10	
T,T,1	Move	T,T,2	1	-2	
T,T,2	Move	T,T,1	1	-2	
T,T,1	Flip	T,T,1	0.5	-1	
T,T,1	Flip	H,T,1	0.5	-1	
T,T,2	Flip	T,T,2	0.5	-1	
T,T,2	Flip	T,H,2	0.5	-1	
T,H,1	Call Fairy	T,H,1	1	-5	
T,H,2	Call Fairy	T,H,2	1	-5	
T,H,1	Move	T,H,2	1	-2	
T,H,2	Move	T,H,1	1	-2	
T,H,1	Flip	T,H,1	0.5	-1	
T,H,1	Flip	H,H,1	0.5	-1	
T,H,2	Flip	T,H,2	0.5	-1	
T,H,2	Flip	T,T,2	0.5	-1	
H,T,1	Call Fairy	H,T,1	1	-5	
H,T,2	Call Fairy	H,T,2	1	-5	
H,T,1	Move	H,T,2	1	-2	
H,T,2	Move	H,T,1	1	-2	
H,T,1	Flip	H,T,1	0.5	-1	
H,T,1	Flip	T,T,1	0.5	-1	
H,T,2	Flip	H,H,2	0.5	-1	
H,T,2	Flip	H,T,2	0.5	-1	

- a.
- b. This is continuing as even when the battery runs out, the robot will continue to operate and as such will not end
- c. Make it such that when the robot runs out of charge, it will stop operating

2.

- a.  $q_{\pi}(s, a) = E_{\pi}[R_{t+1} + \gamma V_{\pi}(S_{t+1}) | S_t = s, A_t = a]$
- b.  $\sum_{s',r} p(s', r | s, a) [r + \gamma v_{\pi}(s')] \text{ as } q_{\pi}(s, a) \sum_a \pi(a | s) = v_{\pi}(s) \text{ and } v_{\pi}(s) = \sum_{s',r} p(s', r | s, a) [r + \gamma v_{\pi}(s')] \sum_a \pi(a | s)$

3.

- a.  $V(s) = \sum_{s',r} p(s', r | s, a) [r + \gamma v_{\pi}(s')] \sum_a \pi(a | s)$

Starting with s=Z

Therefore  $\sum_{s',r} p(s', r | Z, a) [r + \gamma v_{\pi}(s')] \sum_a \pi(a | Z) = V(Z)$

$= 0.5 * 0.8V(Z) + 0.5 * 0.8V(Z)$

$V(Z) = 0.8V(Z)$

$V(Z) = 0$

s=Y

$\sum_a \pi(a | Y) \sum_{s',r} p(s', r | Y, a) [r + \gamma v_{\pi}(s')] = V(Y)$

$$= 0.5 \cdot 1 + 0.5 \cdot 5$$

$$V(Y) = 3$$

$$s = X$$

$$\sum_a \pi(a|X) \sum_{s',r} p(s', r|X, a) [r + \gamma v_\pi(s')] = V(X)$$

$$= 0.5 (0.2 (25))$$

$$= 2.5$$

$$s = W$$

$$\sum_a \pi(a|W) \sum_{s',r} p(s', r|W, a) [r + \gamma v_\pi(s')] = V(W)$$

$$= 0.5 \cdot (3) + 0.5 (0.5 \cdot (4 \cdot 0.8 \cdot 2.5) + 0.5 (2 + 0.8 \cdot 3))$$

$$= 4.1$$

- b. We can change the policy such that  $\pi(a|s) = \pi(b|s) = 0.85$ . This is better than the above policy because if we redo the calculations above, we can get the following values:

$V(Z)$  will remain zero not matter the policy

$$V(Y) = 0.85 \cdot 1 + 0.85 \cdot 5 = 5.1 \text{ which is greater than } 3$$

$$V(X) = 0.85 (0.2 (25)) = 4.25 \text{ which is greater than } 2.5$$

$$V(W) = 0.85 \cdot (5.1) + 0.85 (0.5 \cdot (4 + 0.8 \cdot 4.25) + 0.5 (2 + 0.8 \cdot 5.1)) = 10.064$$

This shows that for all values of  $s$ , the new policy gives us state values greater than the older policy

4.

### **$V(W)$**

$$\text{Episode 1 Val: } 0 + 10 + 0 = 10$$

$$\text{Episode 2 Val: } -10 + 0 = -10$$

$$\text{Episode 3 Val: } 6 + 2 = 8$$

$$\text{Episode 4 Value: } 12 + 0 = 12$$

$$\text{Final Val} = (10 + -10 + 8 + 12) / 4 = 5$$

### **$V(Y)$**

$$\text{Episode 1 Val: } 0$$

$$\text{Episode 2 Val: Non-existent}$$

$$\text{Episode 3 Val: } 6$$

$$\text{Episode 4 Value: } 12$$

$$\text{Final Val} = (12 + 6 + 0) / 3 = 6$$

### **$V(X)$**

$$\text{Episode 1 Val: } 0 + 10 = 10$$

$$\text{Episode 2 Val: } 0$$

Episode 3 Val: Non-existent  
Episode 4 Value: Nonexistent  
Final Val=  $(10 + 0) / 2 = 5$

**V(Z) = 0** as it is the terminal state

5.

a.  $Q(S, A) += \alpha [R + \gamma \max_a Q(S', A) - Q(S, A)]$   
 $= 0.5[4 + 1 \cdot 16 - 0]$   
 $= 0.5 \cdot 20$   
 $Q(W, a) = 10$

b.  $Q(S, A) += \alpha [R + \gamma Q(S', A') - Q(S, A)]$   
 $= 0.5[4 + 1 \cdot 8 - 0]$   
 $= 0.5 \cdot 12$   
 $Q(W, a) = 6$

6.

- a. We use a greedy policy instead of an  $\epsilon$  greedy policy as we want to display the episode with the maximum possible value. We are also not going to be exploring anymore hence we use a greedy rather than epsilon greedy. It will work with Sarsa as Sarsa only selects the action based off the policy passed in so it can work with a standard greedy policy.
- b. This would work with Monte Carlo control(off policy only) because it will converge to the optimal policy  $\pi^*$  as the policy we learn is deterministic and as long as we estimate Q values for it, we learn a policy greedy to the previous estimated target policy we learn and eventually will converge. This however, would not work with On-Policy Monte Carlo control as that is not necessarily guaranteed to give us an optimal policy.