

Computing Science (CMPUT) 455

Search, Knowledge, and Simulations

Martin Müller

Department of Computing Science
University of Alberta
`mmueller@ualberta.ca`

Fall 2022

455 Today - Lecture 5

Today's Topics:

- More on size of state space, effort of solving a game
- Sequential decision-making

Coursework and New Uploads

- Quiz 3
- Assignment 1
- Reading O'Neil, How algorithms rule our working lives
- Activities Lecture 5
- Python codes `count_dag.py`, `generate_tree.py`,
`generate_tree_test.py`

Activities 5a and 5b

- Compute size of game trees
- When does a tree grow more quickly?
 - When increasing b ?
 - When increasing d ?
- Compare two cases

Size and Structure of State Space for Games

Review: State Space vs Playing and Solving a Game

- What is the complexity of solving, or playing well, in a game?
- Depends on many factors:
 - Branching factor
 - Depth
 - Existence of a simple strategy
 - Existence of a mathematical theory
 - Having master players, master games, books to learn from
 - Having good heuristics
 - ...

Review: State Space vs Playing and Solving a Game

Simple measures of complexity of state space

- Size: how many states?
- Structure: tree, DAG, or DCG?
- Branching factor b
- Depth d

Estimating Search Effort

- Assume we need to visit every state in order to solve a game
 - (Later we will see we can do much better)
- How long does it take?
- Main factors:
 - Speed of program
 - Size of state space
- Let's look at 7×7 Go and Go1

7 × 7 Go - Example

- 7 × 7 Go, start on empty board
- Assume we can process 1000 states/second
- Assume simplest tree model, $b = 49$
- What depth d can we reach in which time?
- Can we explore the whole state space?

Brute Force Search for 7×7 Go

Table: Each row shows the estimated *additional effort* to search one level deeper

Depth	New states	Added search time
0	1	1ms
1	49	50ms
2	49^2	2.4s
3	49^3	2 min
4	49^4	1.6 hrs
5	49^5	3.2 days
6	49^6	160 days
7	49^7	21.5 years
...

Fighting Exponential Growth

- We cannot even search 7 moves deep with Go1
- To solve the game we need to see to the end
- This can be over 30 moves deep even for this small board
- Time limits in practice
 - 5 sec - 5 min per move in a tournament game
 - Maybe a few months to solve a game
- How can we succeed?
 - Increase speed of program (Lecture 6)
 - Decrease branching factor b (now)
 - Decrease depth d (later)

Decrease Branching Factor b

- Branching factor = growth of number of states per level
- How to decrease?
 - Reduce number of moves (but how?)
 - Use DAG instead of tree
 - Better search algorithms (e.g. alphabeta search)


Decrease Branching Factor b

- Branching factor = growth of number of states per level
- How to decrease?
 - Reduce number of moves (but how?)
 - Use DAG instead of tree
 - Better search algorithms (e.g. alphabeta search)
- **Question:** How would using a DAG help decrease the branching factor?

Decrease Branching Factor b

- Branching factor = growth of number of states per level
- How to decrease?
 - Reduce number of moves (but how?)
 - Use DAG instead of tree
 - Better search algorithms (e.g. alphabeta search)
- **Question:** How would using a DAG help decrease the branching factor?
- First try: take symmetry into account

Example - Use Symmetry in TicTacToe




- At root: Only 3 of 9 total moves are different
 - Corner (A), Edge (B), Center (C)
- All 6 other moves lead to a symmetric position, same result as A or B
- Symmetries at tree level 1:
 - After corner or edge move: 5 distinct cases
 - After center move: 2 distinct cases
- Limitation: most symmetries broken after few moves

Symmetry

- Typical example to reduce state space by symmetry
- Good reduction at depth 1 or 2
- Then *symmetry breaks*
- Almost no reduction deeper in the tree
- Reduction of whole state space is limited to some constant factor
 - Less than 8 in Go

DAG (Directed Acyclic Graph)



- Idea: single node for all *equivalent* states
- Different paths to same node
- Can lead to huge reduction in state space
- Why?
 - The whole *subtree* below is no longer duplicated
 - Can happen throughout the whole game

Tree vs DAG

- Tree model
 - Each action leads to a new node
- DAG model
 - All *equivalent* states represented by a single node
- Reduction in size of state space
 - Can be many orders of magnitude
 - Examples: Activities 5c - 5e
- Advantages of DAG model:
- Avoid redundant computations
 - No copied subtrees or sub-DAGs
- Share results of analysis - compute once, re-use often

Limitations of DAG Model


- Main problems:
 - Need memory to store and recognize equivalent states
 - Some algorithms designed only for trees, not for DAGs
- Example: propagating information up towards root
 - Only one path up in tree - efficient
 - Many paths in DAG - many ancestors

More Limitations of DAG Model

Limitation: states with different *history*

- Cannot always merge into one node, not always equivalent
- Example: simple ko - is capture allowed?
- Board looks the same but moves are different
- Can you still do something? Yes. One of Martin's students wrote a whole PhD thesis on such questions (Kishimoto 2005)

Counting States in a DAG



- Simplified GoMoku example, also counts illegal states
- Depth 0: 0 black, 0 white stones
- Depth 1: 1 black, 0 white stones
- Depth 2: 1 black, 1 white stones
- Depth 3: 2 black, 1 white stones
- Depth d : $\lceil d/2 \rceil$ black stones and $\lfloor d/2 \rfloor$ white stones
- How many ways to put that many stones on a board with 49 points?

Counting States in a DAG (continued)

- Example:
- Board with 49 squares
- How many different ways to place 5 black stones?
- Answer: $\binom{49}{5} = 1,906,884$
- Need to review the math background? See e.g.
<https://en.wikipedia.org/wiki/Combination>

Counting States in a DAG (continued)

- How many different ways to place 5 black stones and 3 white stones?
- Answer: $\binom{49}{5} \times \binom{44}{3} = 1,906,884 \times 13,244 \approx 25.2$ billion
- Why?
 - $\binom{49}{5}$ ways to place 5 black stones
 - $49 - 5 = 44$ empty points remaining
 - $\binom{44}{3}$ ways to place the 3 white stones there
 - Each different choice for either black or white leads to a different position, so multiply

Activity 5d: Count States in Tic-Tac-Toe DAG

- TicTacToe board has 9 squares
- Ignore symmetry, early wins for now
- At level d : $\lceil d/2 \rceil$ X's and $\lfloor d/2 \rfloor$ O's
- Compute number of positions with 0, 1, 2, 3, ..., 9 stones

0 stones	0 X, 0 O	1 position
1 stone	1 X, 0 O	9 positions
2 stones	1 X, 1 O	? positions
3 stones	2 X, 1 O	? positions
4 stones	2 X, 2 O	? positions
...		? positions
9 stones	5 X, 4 O	? positions

Hint: Python code in `count_dag.py` is useful...

b^d Model and DAG

- Counted the number of states in a DAG at each depth d
- What about the branching factor?
- No longer a constant b for whole DAG
- Different *effective branching factors* b_d depending for each depth (or level) d
- Can compute b_d at depth d as:

$$b_d = \frac{\#nodes-at-depth\ d + 1}{\#nodes-at-depth\ d}$$

- Activity 5e: compute the effective branching factor for the TicTacToe DAG

Computing Size of State Space in DAG from Branching Factors

- Given branching factors, how many nodes in a DAG?
- 1 root node at depth 0
- b_0 children of root $\longrightarrow b_0$ total nodes at depth 1
- Each child has b_1 new children
 \longrightarrow total $b_0 \times b_1$ nodes at depth 2
- Depth n :
 $b_0 \times b_1 \times \dots \times b_{n-1}$ nodes

Computing Size from Branching Factors (continued)

- Total nodes up to depth d
 - 1
 - + b_0
 - + $b_0 \times b_1$
 - + ...
 - + $b_0 \times b_1 \times \dots \times b_{d-1}$
- In general:
no nice closed-form solution for this sum
- In practice:
estimate branching factors b_i
 - Use search or sampling
 - Difficult or impossible to compute them exactly for large games

b^d Model vs Reality: Some Case Studies (1)

- How realistic is the b^d model for size of state space?
- We'll look at some popular games
 - Go, TicTacToe:
 - Roughly, $b_n \approx b_0 - n$
 - Why? One less empty square with each move
 - One less possibility for next move
- Not exact:
 - Ignores illegal move rules
 - Ignores games that end earlier
- Setting $b_n = b_0 - n$ gives $b_0!$ leaf nodes
 - Earlier TicTacToe example: $9 \times 8 \times \dots 1$

b^d Model vs New estimate: 7×7 Go

- 7×7 Go estimate from Lecture 4:
- 25 moves on average during a game, game length about 30 moves
- Rough b^d estimate $25^{30} \approx 10^{42}$
- New model:
- $b_0 = 49$, and $b_n \approx b_0 - n$
- Stop game at $n = 30$
- $49 \times 48 \times \dots \times (49 - 30) = 49!/18! \approx 10^{47}$
- Still ignores captures, ko, different game lengths,...

b^d Model vs Reality: Checkers




Image source:

<https://en.wikipedia.org>

- Checkers: complicated b
- Beginning: many pieces blocked
- Pieces unblocked: b increases
- Forced captures: $b = 1$
- When pieces get captured, b decreases again
- When checkers become kings, b strongly increases
- Estimated average over typical game: $b \approx 2.8$
- Length of game d varies wildly

b^d Model vs Reality: Chess




Image source:

<https://en.wikipedia.org>

- Chess: also complicated
- Pieces such as queens can have many moves, but may be blocked
- King in check: often only few legal moves
- When pieces get captured, b decreases
- Estimated average over typical game: $b \approx 35$
- Length of game d varies wildly

b^d Model vs Reality: Shogi




Image source:

<https://en.wikipedia.org>

- Shogi, Japanese chess
- Similar to chess, plus:
- Captured opponent pieces can be *reused* for yourself in a future move
- With captures, b *increases*
- Estimated average over typical game: $b \approx 92$
- b can be several hundred in endgame with many captured pieces available for “dropping” back on board

Complexity of Popular Games

- Big table in
`https://en.wikipedia.org/wiki/Game_complexity`
- Different measures of complexity
- Complexity also depends strongly on size of board
- In Go, the theoretical complexity is much higher
 - Main reason: capture, play again on same point
 - Example: Go player fills eyes, games last VERY long
 - Game ends only if **all** moves cause full-board repetition

Example: Solving 2×2 Go, $1 \times n$ Go

- 2×2 Example from John Tromp,
`https://tromp.github.io/java/go/twoxtwo.html`
- Naive brute force minimax search: trillions of nodes
- Alphabeta with bad move ordering:
 - 19,397,529 nodes, max. depth 58
- Alphabeta with good move ordering:
 - 1,446 nodes, max. depth 22
- Solving $1 \times n$ Go: *Exploring Positional Linear Go*
paper by Noah Weninger (UofA undergrad!) and Ryan
Hayward (UofA prof) - see resources

Sequential Decision-Making

Sequential Decision-Making

Topics:

- From decision making to sequential decision making
- Notation for action sequences
- View game tree as tree of move sequences

From State Space to Decision-Making

- We studied state spaces in some detail
- Now, how do we find good actions (moves) in a state?
- In general, looking at the current state is not enough
- We need to look ahead to future states in order to make a good decision now
- We need to consider **sequences** of actions, until we reach a terminal state
- In games, each sequence is one possible way of playing the game

Making Complex Decisions

- How to make good decisions?
- Consider many alternatives
- Consider short-term and long-term consequences
- Evaluate different options and choose the best-looking one
- Understanding and comparing sequences of actions is the main step in making such decisions

Making Sequential Decisions

Very general model:

Loop:

- Get current state of world
- Analyze it
- Select an action
- Observe the world's response
- If not done: go back to start of loop

Practically important question:

- Can we do this in a *simulation model* as opposed to the real world?

Single Agent Example: Path Planning

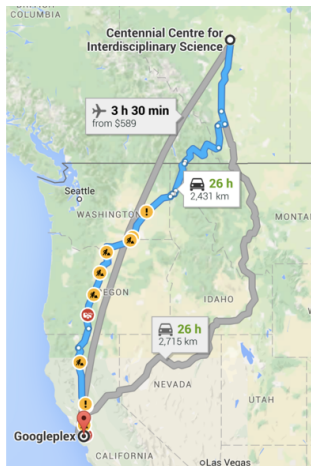


Image source: Googlemaps

- Task: start from here, visit Google headquarters
- First decision: fly, drive, take the bus, or walk?
- If drive or walk: each street corner is a decision point
- Need a long sequence of decisions to arrive at destination
- Is it *optimal*? Is it *good enough*?
- Tradeoffs: Speed vs cost vs scenery vs construction sites ...

Formal Framework

- Sequence of states and actions
- Start state s_0
- Action a_i leads to next state, s_{i+1}
- Keep going until reach a terminal state s_n
- Sequence $(s_0, a_0, s_1, a_1, \dots, s_n)$
- Sometimes we only write the actions (a_0, a_1, \dots, a_n)
 - Example: games where states are determined from game rules and actions

Formal Framework with Rewards

- Formal framework can also include rewards (or costs)
- Simple case (most games we consider):
single reward r at end
- General case: reward r_i after each action a_i
 - Write rewards as part of sequence:
 - $(s_0, a_0, r_1, s_1, a_1, r_2, \dots, r_n, s_n)$

Partial Sequences

- Full sequence goes all the way to terminal state
- *Partial sequence* can stop after any number of actions
- Two full sequences always share a *common prefix*
- In worst case, it might only contain the start state
 - Example - Go game
 - $(s_0, \text{Black B3}, s_1, \text{White A2}, s_{2a}, \text{Black D4})$
 - $(s_0, \text{Black B3}, s_1, \text{White A4}, s_{2b}, \text{Black D4})$
 - Common prefix $(s_0, \text{Black B3}, s_1)$

Re-Interpreting the Tree and DAG Models

- Our model so far:
- State space as a graph
- Nodes are states, edges are actions
- Tree and DAG are special cases of graphs
- New view:
we can view the same trees and DAGs
as a way to organize all action sequences

Organizing Sequences in Trees

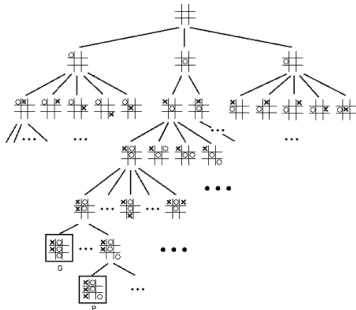



Image source: <http://web.emn.fr>

- Consider the (huge) set of all possible state-action sequences
- Organize them such that:
- Any two sequences share their *longest common prefix*
- Branch as soon as they differ
- Result: we get exactly the tree representation of the state space

Organizing Sequences in a DAG



- Similarly, we can relate sequences to the DAG model
- Start with sequences-as-tree model
- Then, merge two different sequences when they both reach equivalent states
- Result: we get exactly the DAG representation

Summary

- Looked at more details and examples of state spaces
- Estimating size of state space in DAG model
- Theory vs reality: state space of some popular games
- Sequential decision-making model
- Relation between tree and DAG models, sequences of decisions

Preview - What's Next?

- Lecture 6: profiling and optimization
- Topics for next few weeks:
 - Do we need to look at *all* possible sequences to make a decision?
 - How do we use the tree or DAG structure?
 - Algorithms for decision-making in games on tree and DAG structures
 - How do we use heuristics?
 - Simulation using random sequences of actions