

# Computing Science (CMPUT) 455

## Search, Knowledge, and Simulations

Martin Müller

Department of Computing Science  
University of Alberta  
`mmueller@ualberta.ca`

Fall 2022

# 455 Today - Lecture 6

---

## Today's Topics:

- Go rules revisited - more details
- Profiling Python 3 code
- Improving the performance of our Go code

# Coursework and Uploads

---

- Quiz 3 due tonight
- Assignment 1
  - Assignment 1 was due yesterday
  - Feedback from TA via email by end of today
  - Late/second submission deadline **Wednesday 11:55pm**
    - 20% deduction
    - Second submission *allowed for any reason*
  - Your team's last submission counts for marking
- Activities Lecture 6
  - profiling/optimization exercises

# Go Rules Revisited

# Go Rules Revisited

---

- Goal: tidy up some loose ends regarding rules
- Popular variations in rule sets
- Scoring at end of game
- Full repetition rules

# Go Rules So Far

---

- Introduced basic rules
- Showed examples of how to score at the end
- Implemented legal moves and reasonable policy for when to pass at the end in Go1
- Position repetition: implemented only simple ko

# Versions of Go Rules

---

- There are many different versions of Go rules
- All agree on how to handle the vast majority of situations
- Differences in details related to:
  - What is a legal move?
  - When does the game end?
  - How to score the game at the end?
  - How to resolve different opinions about scoring?

# Popular Go Rules

---

- Chinese, Japanese and Korean rules
- Ing, AGA (American Go Association), New Zealand, Tromp-Taylor rules
- Most of these again have different versions and revisions



# Main Differences (1): End of Game

---

- When exactly is the game over?
  - Two or three passes
  - We use **two**

## Main Differences (2): Scoring

---

- How to score at the end?
  - **Area scoring**: count own stones plus surrounded empty points
  - **Territory scoring**: count surrounded points plus captured stones (and prisoners)
  - All rules: add komi to score
  - We use **area scoring**
  - Easier to implement and play correctly at the end
  - In territory scoring, playing inside your surrounded areas costs points




## Main Differences (3): Which Moves are Legal?

---

- Differences regarding suicide and repetition
- Most rules forbid suicide (we do too)
  - Exceptions: e.g. Tromp-Taylor rules
- Repetition: basic ko vs full board repetition
  - Our programs only recognize basic ko

## Review: Repetition Rules - Basic Ko


---



- From top to middle picture: White can capture one black stone by playing A
- From middle to bottom picture: Now if Black captures back one white stone...
- The position would repeat, infinite loop
- This is called a (basic) ko.
- Go rules forbid such repetition

# Repetition - Longer Loop

---










A four move loop in Go. Black passes on move 3.

- Example of a longer repetition loop
- This really happens in games between weaker Go programs
- If White tries to play move 4 in the corner, it repeats the position from four moves ago
- If both continue like this, infinite loop
- Go1 does not recognize or prevent such repetition

# Repetition - Triple Ko

---



A triple ko leading to a six move long loop in Go.

# Full Board Repetition

---

- Many rule versions forbid that the same board position is repeated in a game
- In the examples, the last, loop-closing move is illegal
- Such rules are often called *superko* rules
- They handle complex loops and situations with multiple active ko

# Positional vs Situational Superko

---

- Superko idea: do not repeat the same board position
- What exactly is “the same”?
- Two main answers:
- Positional superko (PSK)
  - Ignore whose turn it is, only compare board
- Situational superko (SSK)
  - Compare whose turn it is as well as board
- Even more details: how do pass moves affect the repetition ban?



# Detecting Superko Repetition

---

- Simplest but slowest:
  - Compare against all previous positions
  - Much too slow in practice
- One solution: use hashing to detect potential repetition
- Simple, effective trick (not complete solution):  
check if a move has ever been played before
- No details now, some later in search chapter

# Profiling and Code Optimization

# Profiling and Code Optimization

---

- Our `Go0` and `Go1` Python sample codes are very slow
- They were written for simplicity, not speed
- This is usually a good first approach - see quotes next slide
- Optimization is very important in search, but it can wait a bit
- We can optimize *if* and when we need it
- First, look where the time is spent
- **Profiling** is an easy way to check this

# Go Code Now vs Then

---

- Note: these slides reflect my real experience from a few years ago
- At that time, we used an earlier version of our code base
- Some details are different now
- The “big picture” and the main messages are still valid

## Some Famous Quotes

---

Fred Brooks, The Mythical Man-Month (1975)

*The management question, therefore, is not whether to build a pilot system and throw it away. You will do that. [...]*

*Hence plan to throw one away; you will, anyhow.*

Don Knuth, Structured Programming with go to Statements (1974)

*We should forget about small efficiencies,  
say about 97% of the time:  
premature optimization is the root of all evil.*

# Limits of Optimization

---

- There is often an (approximate) 80-20 rule:  
80% of the improvement can come  
from 20% of the code
- With search, it can be even higher
- However, **Amdahl's law** limits the amount of speedup

# Limits of Optimization

---

- There is often an (approximate) 80-20 rule:  
80% of the improvement can come  
from 20% of the code
- With search, it can be even higher
- However, **Amdahl's law** limits the amount of speedup

## Example

- Assume a program spends 80% of its time in one function
- We manage to speed this function up 100x
- **Question:** How much is the overall speedup?

# Limits of Optimization

---

- There is often an (approximate) 80-20 rule:  
80% of the improvement can come  
from 20% of the code
- With search, it can be even higher
- However, **Amdahl's law** limits the amount of speedup

## Example

- Assume a program spends 80% of its time in one function
- We manage to speed this function up 100x
- **Question:** How much is the overall speedup?
  - Less than 5x



# Amdahl's Law




Image source:

[https://en.wikipedia.org/wiki/Amdahl's\\_Law](https://en.wikipedia.org/wiki/Amdahl's_Law)

Amdahl's Law

- Amdahl's Law (1967)
- How does speeding up one part of program speed up the whole?
- Often used for parallel programming
- Main idea: the parts of the program that are not optimized limit the overall speedup

# Amdahl's Law - Formula

---

- $p$  = percentage of program that is speeded up
- $s$  = speedup for that part
- Runtime before optimization: 1
- Runtime after optimization:  $(1 - p) + p/s$
- Speedup limit for the whole program:
  - $\text{limit} = \frac{\text{runtime before}}{\text{runtime after}} = \frac{1}{(1-p)+p/s}$
- Simplified version: assume  $s$  very large, then  $p/s$  is very small, ignore ...
  - $\text{limit} \approx \frac{1}{1-p}$

# Amdahl's Law - Example Revisited

---

- 80% of program speeded up, so  $p = 0.8$
- $s = 100$  speedup for the optimized function
- Speedup limit for the whole program:

- $\text{limit} = \frac{1}{(1-p)+p/s} = \frac{1}{(1-0.8)+0.8/100} \approx 4.81$

- Simplified version:

- $\text{limit} \approx \frac{1}{1-p} = 5$

# Profiling

---

- Define a test that runs your program with a typical workload
- Run it with a special program called **profiler**
- Profiler tells you details of the program execution
- Profilers can be on the function level or instruction level
- How often was piece of code executed?
- How long did it take?
- Possibly, lower level details such as cache misses

# Simple Profiling in Python with cProfile - Code

---

See code `profile_Go1.py` in directory `go0and1`

```
import cProfile
from Go1 import Go1
...
def play_moves():
    """
    play 100 random games of 100 moves each
    for profiling.
    """
    ...
cProfile.run("play_moves()")
```

# Simple Profiling in Python

---

- See code `profile_Go1.py`
- Try it out with
- `./profile_Go1.py > profile.txt`
- `sort -k 2 -r profile.txt`
- This sorts by total time per function
- Try other options for `-k` to sort by other criteria
- Example: `sort -k 1 -r profile.txt`

# Ways of Profiling in Python

---

- `cProfile` is a built-in module, no need to install anything
- Downside: overhead of profiling is also measured
- More advanced profilers are available for download:
  - `Profilehooks`
  - `pycallgraph`

See profiling on our Python language page

# Speeding Up Go1

---

- Go1 is slow
- For search and simulation, speed is very important
- How to improve the code?
- Both low-level optimizations and better algorithms help
- Case study: a series of improvements to Go1
  - Result: Go2 - same algorithm as Go1 but faster



# An Iterative Optimization Procedure

---

- First, pick a test to measure the speed
- Here: play 100 games on  $7 \times 7$  board
- Repeat:
  - Run test games with profiler
  - Identify the most expensive functions
  - Try to improve them by optimization or better algorithms

# Profiling Go1

---

- Profile with `cProfile`
- Total time: 6.2 seconds
- Worst 5 individual functions listed below  
(all in `board.py`)

Calls	Time	Name
561025	1.960	<code>neighbors_of_color</code>
2287541	0.680	<code>get_color</code>
610480	0.679	<code>_neighbors</code>
43441	0.662	<code>_block_of</code>
18268	0.405	<code>play_move</code>

# Profiling Go1

---

- Also look at cumulative time
- Function itself plus other functions it calls
- Sort by column 4:  

```
sort -k 4 -r profile.txt
```
- Some interesting functions listed below  
(all in `board.py`)

Calls	Cumulative Time	Name
10974	4.429	<code>is_legal</code>
25584	3.566	<code>_detect_and_process_capture</code>
43441	3.368	<code>_block_of</code>
561025	3.351	<code>neighbors_of_color</code>
43441	1.359	<code>_has_liberty</code>

# Strategies for Optimization

---

- Best: avoid calling a function
- Second best: speed up a function, avoid unneeded computation
- Here: detecting captures is most expensive

## Read the Code

---

- Start by reading the expensive code carefully
- Can we avoid unneeded computation?
- Here: read `_has_liberty`, `neighbors_of_color`

```
def _has_liberty(self, block):  
    for stone in where1d(block):  
        empty_nbs = self.neighbors_of_color(  
            stone, EMPTY)  
        if empty_nbs:  
            return True  
    return False
```

## Read the Code

---

```
def neighbors_of_color(self, point, color):  
    nbc = []  
    for nb in self._neighbors(point):  
        if self.get_color(nb) == color:  
            nbc.append(nb)  
    return nbc
```

- We do not need to compute the whole list
- Stop if we find one liberty
- `neighbors_of_color` is still used in other places
- Add a function that is optimized for our task

## New Version

---

```
def find_neighbor_of_color(self, point, color):
    for nb in self._neighbors(point):
        if self.get_color(nb) == color:
            return nb
    return None

def _has_liberty(self, block):
    for stone in where1d(block):
        if self.find_neighbor_of_color(stone, EMPTY):
            return True
    return False
```

## Profiling Again

---

- Total time reduced from 6.2 to 6 seconds
- Reduction in `_has_liberty` by calling cheaper `find_neighbor_of_color` instead of `neighbors_of_color`
- Nice improvement for a little work, but not a huge win
- Can we avoid the many floodfills altogether?
- We do the floodfill for each neighbor of a stone
- We only need to know “does block have at least one liberty”?
- Can we check that more effectively?



# Optimizing Floodfill

---

- We can store such a liberty for each stone  $s$
- In the code: `liberty_of[s]`
- Check capture: just check if board at location `liberty_of[s]` is still empty
- If yes, no floodfill is needed (why?)
- If no, we just played there
  - Do floodfill to try to find a *different* liberty for  $s$
  - If success: update `liberty_of[s]`
  - If fail: yes it is a capture

## Result, and More Floodfill Optimization

---

- Total time reduced from 6 to 4.4 seconds
- Success!
- Next: try to reduce calls to expensive floodfill functions
- Idea: instead of always computing a block:
- First check the 4 neighbors of the stone if there is a liberty there
- Result: Total time reduced from 4.4 to 3.7 seconds
- Cost: more complex code, adds special case

## Profiling Again

---

Calls	Time	Name
66323	0.669	find_neighbor_of_color
18645	0.396	play_move
32369	0.367	_is_surrounded
264389	0.321	_neighbors
147455	0.294	neighbors_of_color
828018	0.257	get_color

## Profiling Again

---

Calls	Time	Name
66323	0.669	find_neighbor_of_color
18645	0.396	play_move
32369	0.367	_is_surrounded
264389	0.321	<b>_neighbors</b>
147455	0.294	neighbors_of_color
828018	0.257	get_color

## Optimizing Neighbors, First Try

---

```
def _neighbors(self, point):  
    return [point-1, point+1,  
            point-self.NS, point+self.NS]
```

- Called often: compute list of neighbors of a point
- Each call creates a new list
- Some neighbors are off the board (state `BORDER`), causing more tests in code
- Precompute a `neighbors` array for each point
- Include only on-board neighbors
- Result: EPIC FAIL, runtime over 11 seconds
- Why? board is copied and neighbors array recomputed over 11000 times

## Optimizing is\_legal

---

```
def is_legal(self, point, color):  
    board_copy = self.copy()  
    legal = board_copy.play_move(point, color)  
    return legal
```

- This function is the reason for FAIL with previous optimization
- Slow: copy the board, then try to play the candidate move to see if it is legal
- Solution: Implement `is_legal` without `play_move`
- Success! Total time reduced from 4.4 to 2.5 seconds
- Cost: increased code complexity, some redundancy in `is_legal` and `play_move`

## Details

---

Calls	Time	Name
51038	0.528	find_neighbor_of_color
75984	0.288	neighbors_of_color
21163	0.227	_is_surrounded
166427	0.207	_neighbors
495786	0.181	get_color
7418	0.145	play_move
{prev: 18645	0.396	play_move}

- `play_move` calls: less than half as many
- Many other function calls also significantly reduced

## Optimizing Neighbors, Second Try

---

- Now we are no longer copying the board at each legal move check
- Now the neighbors optimization works beautifully
- Result: Total time reduced from 2.5 to 2 seconds
- Success!
- There are more opportunities to optimize but Martin stopped here



# Summary (1)

---

- Discussed profiling and optimization
- Some concrete case studies
- Overall about 3x faster now, from 6 to 2 seconds on test
- Strategies:
  - Save computation
  - Precompute
  - Compute data incrementally when there are only small changes
  - Catch and handle frequent simple cases early

## Summary (2)

---

- Very few optimizations are win-win. The speed often comes at the cost of code complexity
- Remember Knuth:  
premature optimization is the root of all evil