

Computing Science (CMPUT) 455

Search, Knowledge, and Simulations

Martin Müller

Department of Computing Science
University of Alberta
`mmueller@ualberta.ca`

Fall 2022

455 Today - Lecture 16

- Today: Selective search and Monte Carlo Tree Search (MCTS)
- Finish discussion of UCB from Lecture 15
- Comparison and overview:
 - Exact search
 - Selective search
 - Simulations
- Monte Carlo Tree Search framework
- UCT algorithm
- Enhancements of MCTS

Coursework

- Work on Assignment 3
- Reading: Pedro Domingos, A Few Useful Things to Know about Machine Learning
- Quiz 9 Monte Carlo Tree Search. Double length
- Activities

Exact Search, Selective Search, and Simulations

- Big-picture overview of algorithms so far
- For each method, focus on three questions:
 1. Which parts of the game tree does it visit?
 2. How does it back-up results to the root of the tree?
 3. Exact or selective?


Review - (b,d) Tree Model and Solving a Game

- Search space in (b,d) tree model:
- Branching factor b , depth d
- Alternating min and max levels in tree
- b^d leaf nodes
- $(b^{d+1} - 1)/(b - 1) \approx (b * b^d)/(b - 1) \approx b^d$ nodes in whole tree
- Size of proof tree in best case: very roughly $b^{d/2}$
- Minimum amount of search to solve a game

Naive Minimax (or Negamax) - Exact Solver

1. Which parts of the game tree does it visit?
 - Explores the full game tree
 - All children searched in each node
2. How does it back-up results to the root of the tree?
 - Minimax:
Minimum over children at min nodes,
maximum at max nodes
 - Negamax is a different but equivalent formulation, same result
3. Exact or selective?
 - Exact
 - Terminal nodes are true end-of-game
 - Uses only exact scores at terminal nodes for evaluation
 - Result is proven correct

Naive Minimax - Exact Solver




Efficient Minimax (or Negamax) - Boolean Minimax, Alphabeta

1. Which parts of the game tree does it visit?
 - Some parts of tree may be cut by exact pruning rules
 - Best case: Visit only 1 child for winner
 - Needs to try all moves for loser
2. How does it back-up results to the root of the tree?
 - Minimax
 - For alphabeta, some back-up values are “good-enough” upper or lower bounds, not exact values
3. Exact or selective?
 - Exact.

Efficient Minimax (or Negamax) - Boolean Minimax, AlphaBeta


Best case:
Search
proof tree
first



Depth-limited Alphabeta Search

1. Which parts of the game tree does it visit?
 - As in alphabeta, but only up to depth limit
2. How does it back-up results to the root of the tree?
 - Min and max
3. Exact or selective?
 - Selective
 - Heuristic evaluation at terminal nodes
 - Search process is exact, but evaluation of leaves is not
 - Source of error: heuristic evaluation in leaf nodes

Depth-limited Alphabeta Search



Selective Alphabeta Search with Fixed Time or Node Budget

- For many games even $O(b^{d/2})$ nodes for best-case proof is far too large
- In practice: fixed time or node limit for the search (e.g. 30 seconds, or 10^{12} nodes)
- What can we search within that budget?
- First answer was depth-limited search: reduce d until search fits within budget
- *New: Second answer - selective search:* reduce both b and d until search fits within budget


Selective Alphabeta Search - Methods

- How to do selective search?
- Search “interesting” moves much deeper than others
- Choice 1: Prune moves by using selective minimax algorithms such as ProbCut (Buro) or Nullmove pruning
- Choice 2: Prune moves using knowledge
 - Details: <https://www.chessprogramming.org/Selectivity>
- Choice 3: expand search tree selectively
 - Example: Monte Carlo Tree Search (MCTS)

Selective Alphabet Search

1. Which parts of the game tree does it visit?
 - Does not consider all legal moves in each node
 - Often depth-limited as well
2. How does it back-up results to the root of the tree?
 - Min and max
3. Exact or selective?
 - Selective
 - Heuristic evaluation at terminal nodes
 - Skips some legal moves
 - Source of error: heuristic evaluation in leaf nodes
 - Source of error: may prune the best move from a node

Selective Alphabeta Search



Selective Alphabet Search for Large Problems


- Large problems (chess, checkers, ...)
- Reducing b not enough
- Reduce both b and d - selective search with heuristic evaluation
- Before Monte Carlo, this was the standard approach for most complex games

Selective Alphabet Search for Go?

How about Go?

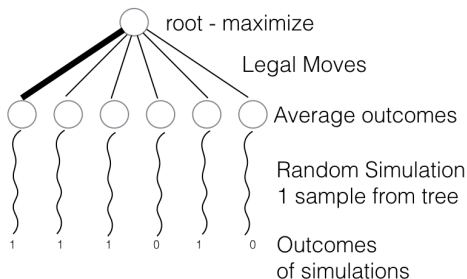
- > 200 moves on average for 19x19 Go
- Usually, only 1-10 of them are good
- Can we reduce b down to this range, without missing important good moves?
- Many attempts failed in the past - too many good moves missed
- MCTS was the first approach that worked well
- Later, strong move selection heuristics based on neural nets also helped a lot
 - Neural nets were not tried much with alphabeta, since MCTS worked so well in Go

Simulation-based Players



- Review: simple simulation-based players (e.g. Go3)
- 1 ply search at the root
- Move selection - simple or UCB
- Simulations - (almost) random, rule-based, or probabilistic
- How do these algorithms compare to selective search?

Simulation-Based Player as Selective Minimax Search




- Extreme case of selective search:
- Simulation-based player with one simulation per move
- Branching factor b at root, complete search
- Branching factor 1 at all later levels...

Simulation-Based Player with Repeated Sampling

Repeated sampling in Simulation Player

- With small number of samples
 - Samples a few moves close to the start
 - does not improve the branching factor lower in the tree
- With large, unlimited number of samples
 - Eventually samples all nodes in the full (b,d) tree infinitely often
 - With selective policy (e.g. patterns, filters), samples some subtree infinitely often

Simulation-based Player - Uniform Random Simulation Policy




Simulation-based Player - Uniform Random Simulation Policy

1. Which parts of the game tree does it visit?
 - Eventually, visits all nodes
2. How does it back-up results to the root of the tree?
 - Max at root only
 - Average over all simulations
3. Exact or selective?
 - Selective
 - Not exact because of averaging instead of minimax
 - Source of error/risk: bias - average may be far from min, max
 - Source of error: variance - large uncertainty with small number of samples

Simulation-based Player - Non-Uniform Simulation Policy

1. Which parts of the game tree does it visit?
 - All, except subtree below moves that are never selected by policy
2. How does it back-up results to the root of the tree?
 - Same as Uniform Random: 1-ply max + average over simulations
3. Exact or selective?
 - Selective
 - Similar to Uniform Random
 - Strength: average over better samples may be closer to min, max
 - Risk: can miss totally by hard-pruning all good moves

Simulation-based Player - Non-Uniform Simulation Policy



- Some nodes in tree may never be sampled:
- If some move on path to node never selected by policy

Simulation-based Player - Simple vs UCB Move Selection


- Move selection:
both simple and UCB behave the same in principle
- Both compute average over all simulations
- Difference: in UCB, average is taken:
 - Over more simulations for good move
 - Over fewer simulations for bad move
 - With a tree, in MCTS with UCT, this will be important
- Another difference:
 - UCB selects most-simulated move (**Why?**)
 - Simple selects move with highest winrate
 - These moves are usually, but not always the same

Monte Carlo Tree Search (MCTS)

Next algorithm: Monte Carlo Tree Search (MCTS)

1. Which parts of the game tree does it visit?
 - Tree search at the start, simulations to finish
2. How does it back-up results to the root of the tree?
 - Weighted averages over children
 - Weight of child = number of simulations for that child
 - Approaches min, max if best child has much higher weight than rest
3. Exact or selective?
 - Selective
 - Much deeper search for moves with better winrates
 - Converges to exact if given enough time to grow whole tree
 - Weighted average

Monte Carlo Tree Search



Monte Carlo Tree Search

- Weakness of simulation-based players so far:
- No tree search after move 1
- Everything from move 2 is random(ized) simulations only

MCTS + UCT approach

- Add selective tree search
- Adapt UCB idea to work in trees - UCT algorithm
- UCT = Upper Confidence bounds on Trees
- Run simulation from leaf of tree for evaluation

Adding a Game Tree to Simulation-Based Player

- First idea: combine what we have:
 - Depth-limited alphabeta
 - Evaluation by simulation
- This fails miserably.
 - Too noisy - need many simulations to get reasonably stable evaluation
 - Too slow - even 1-ply simulation-based player is slow
 - Result: for many years, simulation-based approaches were ignored for games without random element (no dice, no cards)
- Smarter way to combine search and simulation
 - MCTS + UCT

Monte Carlo Tree Search(MCTS) Model




Fig. 2. One iteration of the general MCTS approach.

Image source: Browne et al, A Survey of Monte Carlo Tree Search Methods

Four steps, repeated many times

- Select - traverse existing tree using formula such as UCT to select a child in each node
- Expand: add node(s) to tree
- Simulate: follow randomized policy to end of game
- Backpropagate: update winrates along path to root

Using MCTS to Play Games

- To play one move:
 - Run MCTS search from current state
 - After search: select best move at root, play it
- To play a whole game:
 - Run MCTS every time it is your program's turn
 - May store and re-use parts of the tree from previous search

Monte Carlo Tree Search(MCTS) Model

Algorithm 1 General MCTS approach.

```
function MCTSSEARCH( $s_0$ )  
  create root node  $v_0$  with state  $s_0$   
  while within computational budget do  
     $v_l \leftarrow \text{TREEPOLICY}(v_0)$   
     $\Delta \leftarrow \text{DEFAULTPOLICY}(s(v_l))$   
    BACKUP( $v_l, \Delta$ )  
  return  $a(\text{BESTCHILD}(v_0))$ 
```

Image source: Browne et al, A Survey of Monte Carlo Tree Search Methods

- v_l = leaf node in tree
- Δ = result of simulation
- $a(..)$ = action to move to best child

Monte Carlo Tree Search Example




Image source: David Silver

Monte Carlo Tree Search Example




Image source: David Silver

Monte Carlo Tree Search Example




Image source: David Silver

Monte Carlo Tree Search Example




Image source: David Silver

Monte Carlo Tree Search Example




Image source: David Silver

MCTS Tree Traversal

- Start from root of tree
- Repeat:
 - Go to best child
 - Until reached leaf node in tree
- What is the best child?
- Use a formula to evaluate all children
 - UCT is popular (see next slides)
- Many other more complex formulas are possible
 - Example: add knowledge-based terms

From UCB to UCT

- UCT algorithm by Kocsis and Szepesvari (2006)
- It is still *the* classic algorithm for Monte Carlo Tree Search
- It is not the first child selection algorithm used in MCTS ...
- ... but it is the first based on sound theory
- Worked better in practice than earlier ad hoc algorithms
- Original paper has over 3300 citations -
hugely influential

UCT Algorithm Main Ideas

- Algorithm for child selection in Monte Carlo Tree Search
- Name UCT is often used for MCTS with this algorithm
- Combines tree search with simulations
- Uses results of simulations to guide growth of the game tree
- Uses UCB-like rule to select “best” child of a tree node
- Goal: select a good path in the tree to explore/exploit next
- Grows the tree over time
- Stores winrate statistics in each node, used for child selections

Exploration vs Exploitation

- Like UCB, UCT tries to balance Exploration and Exploitation
- Exploitation: focus on most promising moves
- Exploration: focus on moves where uncertainty about evaluation is high
- Difference: evaluate UCT formula in every node along a path in the search tree

From UCB to UCT

- Review - UCB formula

$$UCB(i) = \hat{\mu}_i + C\sqrt{\frac{\log N}{n_i}}. \quad (1)$$

- UCT is very similar: UCT value of move i from parent p :

$$UCT(i) = \hat{\mu}_i + C\sqrt{\frac{\log n_p}{n_i}}. \quad (2)$$

- Only difference in exploration term
 - UCB: uses global count of all simulations N
 - UCT: uses simulation count of parent n_p
- For root, UCT is identical to UCB
 - N = simulation count of root

MCTS Tree Expansion

- How to grow the tree?
- Simplest case: add one node per iteration
- Add one node from current simulation
- Tree grows very selectively
 - paths with strong moves become much deeper than others
- Expansion is *optional* - if memory fills too quickly:
 - Use an *expansion threshold* t_e
 - Only add a node if the leaf has at least t_e visits
 - Example: Fuego program, default $t_e = 3$

MCTS Simulations

- Run one simulation from the leaf node of tree
- Can use any simulation policy
 - Uniform random, rule-based, or probabilistic
- Result of simulation is win (1) or loss (0)
- Can run more than one simulation from each leaf node
 - Tradeoff between speed and accuracy
 - Tradeoff between time spent in updating tree vs running simulations
 - Example: for Fuego, on some hardware 2 simulations per leaf works better than 1

MCTS Backpropagation - Update Statistics

- Update wins and visit counts along path to root
- Negamax style implementation - flip wins/losses at each step
- `value = 1-value` changes from wins to losses and back

```
def backprop(node, value):  
    while node:  
        node._wins += value  
        node._n_visits += 1  
        value = 1 - value  
        node = node._parent
```

MCTS Move Selection

- Run as many iterations of MCTS as you can
- Then select move to play at root
- How?
- Browne's paper mentions several approaches
- We discuss the main ones

MCTS Move Selection

- Max child: child with highest number of wins
- Robust child: Select the most visited root child. (This is popular)
- Highest winrate
 - Not a good/stable method with MCTS
 - Why not stable: see next slides
- Max-Robust child (see later slide)

Dangers of Selecting Move by Winrate in MCTS

- MCTS usually expands the move with best winrate (exploitation)
- But sometimes, it explores an inferior-looking move
- This can lead to trouble for selecting a move by best winrate
- A move with low simulation count and high uncertainty about its value might get selected
- See example next slide

Dangers of Selecting Move by Winrate in MCTS

- Example: two moves A and B
- A 78 wins / 100 visits, winrate 78%
- B 6 wins / 8 visits, current winrate 75%
- Assume B has higher UCT score, so we explore B
- B gets a win, now has 7 wins / 9 visits, current winrate 77.8%
- Explore B again
- B gets another win, now has 8 wins / 10 visits, current winrate 80%
- Assume we stop search now

Dangers of Selecting Move by Winrate in MCTS

- A 78 wins / 100 visits, winrate 78%
- B 8 wins / 10 visits, **winrate 80%**
- If we select B because of highest winrate:
- High risk of being wrong
- The value of A is much more certain
- The value of B still has much higher variance
- Remember discussion of binomial distribution of simulations
- Probability of error is high

Max-Robust Child: Extending Search

- What if most-simulated move and highest winrate move are different?
- Search may just have found a new best move
 - B is really better than A
- Or B may be a fluke
 - B got some “lucky” wins, but is worse than A in the long run
- Very little evidence to decide which is true
- One solution: extend the search in such cases

Max-Robust Child: Extending Search

- Extending the search can distinguish two cases:
- If B is really good:
 - B will now receive many more simulations soon, stabilize value
- If B's recent wins were a fluke:
 - Its winrate and upper confidence bound will drop quickly with more simulations
- Extending search in this way is called “Max-Robust child” in the paper

Improving MCTS

- Many ways to improve:
- Adding knowledge in tree or in simulation
- Modify in-tree selection
- Modify or replace simulations
- We will discuss several good options when we talk about machine learning and AlphaGo

Summary

- Overview of game tree search and simulation
- Discussed Monte Carlo Tree Search
- After all the preparation, MCTS mostly combines previously discussed concepts
- 4+1 steps of MCTS
 - Repeat: select, (expand), simulate, backpropagate
 - Finally: select move to play